



Big Data, easy as Pi

Phillipe Loher

Outline

- * Background
- * Hadoop
- * Amazon Web Services
- * Word Count example 'hands-on deployment'
- * Pi 'Battle the Greeks' via Monte Carlo simulation

‘Big Data’ buzz

- * Extremely large datasets that can be analyzed to reveal patterns, trends, and associations
- * Affecting every industry:
 - * Cancer/genetics research, climate and weather forecasting, physics simulations, supply chain optimization, food recipe optimization, ...
- * Changes the way experiments are done
- * New algorithms are being formed
- * Will add fuel to artificial intelligence and the semantic web

Traditional SAN based storage example




Image from: http://commons.wikimedia.org/wiki/File:Sch%C3%A9ma_SAN.png

A little background/history

- * Multi-threading (e.g. POSIX threads, OpenMP)
- * Job Scheduling Systems (e.g. PBS)
- * MPI (Message Passing Interface)

What is Hadoop

- * Software framework for distributed storage and distributed processing of Big Data
- * Open source and Java based – inspired by Google's MapReduce & Google File System (proprietary & C++)
- * Created in 2005, Yahoo is one of many prominent contributors
- * 'Hadoop' often refers to the entire Hadoop Ecosystem which includes packages built on top of Hadoop (such as Apache Hive and Pig)

Hadoop disadvantages - not for everyone

- * Not all algorithms can be easily tailored to fit within the framework
- * Jobs don't communicate with one another
- * Finding 'best' MapReduce settings across the nodes for a particular workload is not always obvious
- * Not real time (work in progress) – batch processing
 - * see Google Caffeine - http://www.theregister.co.uk/2010/09/09/google_caffeine_explained/

Hadoop advantages

- * Well-defined framework that makes it easy to start utilizing distributed computing
- * Reliable and fault tolerant
- * Once your algorithm is written, can easily scale up by adding more compute power
- * Bring the computing to the data (not the other way around)
- * Free and Open Source

Hadoop Common Components

- * **Hadoop Common** – common libraries/utilities
- * **Hadoop YARN** – “Yet Another Resource Negotiator”
- * **Hadoop Distributed File System (HDFS)** – a file system designed to run in a reliable and redundant manner across many low commodity machines
- * **Hadoop MapReduce** – software framework that facilitates processing huge amounts of data on a large number of compute nodes, while reliably handling hardware and infrastructure failures
- * The map reduce framework takes advantage of the distributed nature of HDFS (move workload to data)

Hadoop Distributed File System (HDFS)

- * Namenode/Master (meta-data)
- * Datanode (distributed)
- * Large block sizes (64Mb / 128 Mb typical)
- * Each data block replicated multiple times (for reliability and performance)
- * Better for larger files

Word Count

- * **Problem:** mine extremely large text files and count the # of occurrences of each word
- * **Mapper:** write our own leveraging Hadoop Streaming
- * **Reducer:** can write your own or use the built-in aggregator

<http://hadoop.apache.org/docs/current/api/org/apache/hadoop/mapred/lib/aggregate/package-summary.html>

MapReduce Overview

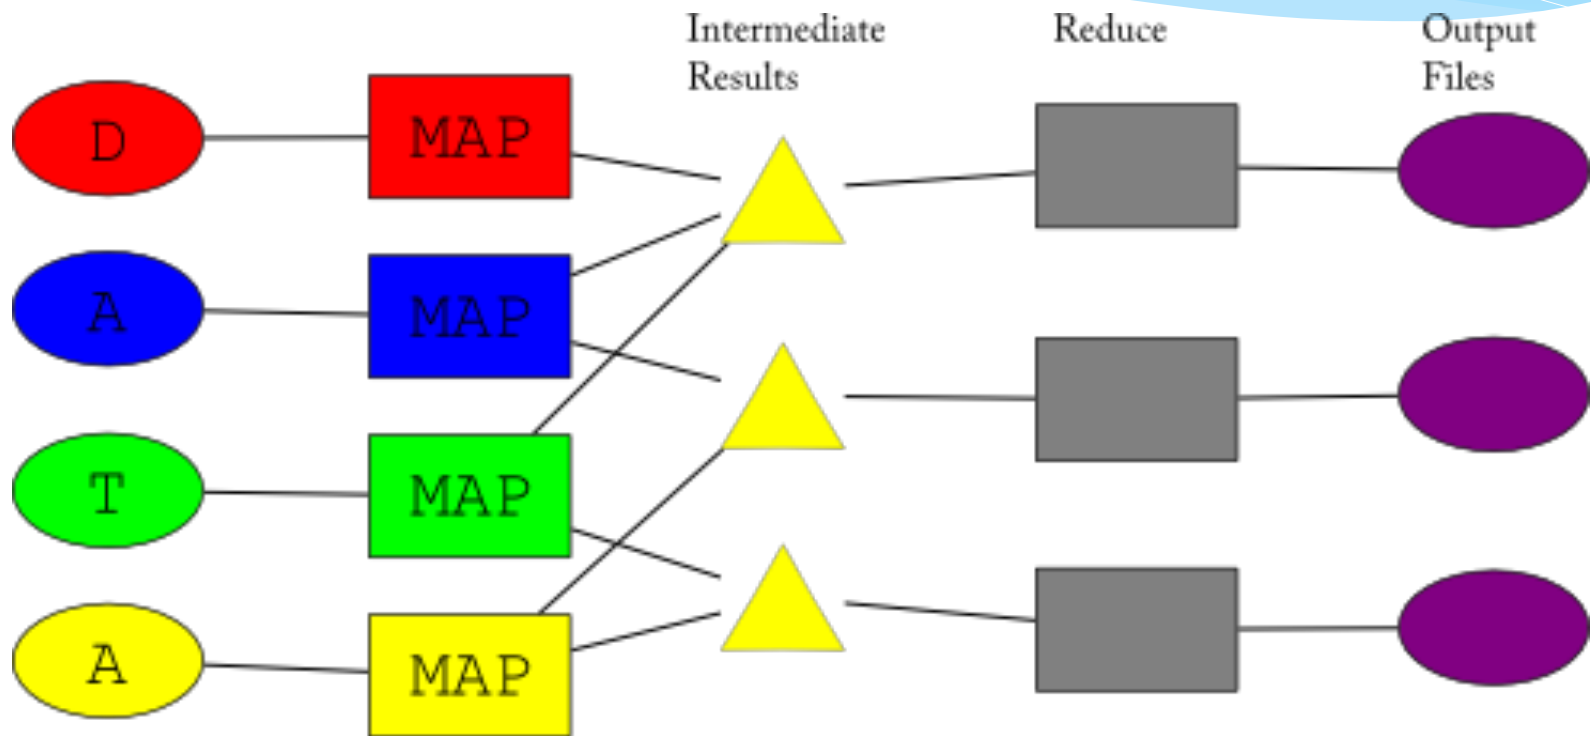


Image from: [http://commons.wikimedia.org/wiki/File:Mapreduce_\(Ville_Tuulos\).png](http://commons.wikimedia.org/wiki/File:Mapreduce_(Ville_Tuulos).png)

Mapper Task

- * Parallel Workhorse
- * Transform input records to intermediate records
- * Records are stored in key/value pairs
- * **Other:**
 - * Specifying a mapper is required
 - * User can 'recommend' how many Mappers will be invoked. Ultimately determined by number of data blocks composing input files
 - * Input files are optional (Pi example at end)

Reducer Task

- * Takes intermediate records and reduces them
- * Again, in key/value pairs
- * **Important:** All matching keys from the Mapper's go to the same Reducer
- * Can't start until all mappers have completed (though transferring/copying can begin)
- * **Other:**
 - * Can set number of reducers when running job
 - * Reducers are optional (mapper-only workloads)

Hadoop Streaming

- * Part of Hadoop that allows writing the Mapper and Reducer in any language you want
- * Uses STDIN/STDOUT
- * Disadvantages: Inter-process overhead and can't override/extend many of the MapReduce constructs

Mapper_Wordcount.pl

```
# read one line at a time
while (my $inputline = <STDIN>)
{
    # split the line into individual words
    my @matches = ($inputline =~ m/[a-zA-Z][a-zA-Z0-9]*/g);
    foreach my $match (@matches)
    {
        # print each word (in lower case)
        printf ("LongValueSum:%s\t1\n", lc ($match));
    }
}
~
```


Reducer_Wordcount.pl

```
# read in <key, value> pair from intermediate results generated by the MAPPER
foreach my $input (<STDIN>)
{
    chomp $input;

    # only get the key, throw away the value since we know it's always 1
    my $key = $input;
    $key =~ s/^LongValueSum://;
    $key =~ s/\t.*//;

    # aggregate the counts for that word (key)
    if (!defined ($aggregate{$key}))
    {
        $aggregate{$key} = 1;
    }
    else
    {
        $aggregate{$key} = $aggregate{$key} + 1;
    }
}

# print the count for each word (key)
foreach my $key (keys %aggregate)
{
    printf ("%s\t%d\n", $key, $aggregate{$key});
}
~
~
```

Resource Management

- * **Jobtracker**

- * Receives the job submission and dispatches the work to all the tasks
- * Scheduling
- * Speculative execution (optional for stragglers)
- * Fault tolerance handling
- * Shuffle & Sort, synchronization

- * **Tasktracker**

- * Sends heart beat messages to Jobtracker
- * Responsible for running the user's code
- * Sends machine availability info (such as free slot availability)
- * If using HDFS, runs alongside datanode daemon

- * **YARN**


- * New in Hadoop 2
- * ResourceManager & NodeManager

Amazon AWS


- * **Amazon Web Services** – collection of remote computing services that form a cloud computing environment.
- * **Amazon Simple Storage Service (S3)** – secure, durable, and scalable cloud-based storage
- * **Amazon Elastic Compute Cloud (EC2)** – virtual computing environments
 - * Different size (memory, cpu, storage, network capabilities) instances
 - * Can pre-configure using Amazon Machine Images (AMIs)
- * **Amazon Elastic Map Reduce (EMR)** – allows for quick deployment of Hadoop using the cloud



Hands-on *AWS* / Hadoop demo



Pi



- * Using the Pythagorean Theorem and a 96-sided polygon, Archimedes of Syracuse (born 287 BC), approximated Pi to be between ~ 3.140845 and ~ 3.142857
- * Pi is the 16th letter in the Greek alphabet
- * Using distributed processing using MapReduce, can you do better than ancient math wizard Archimedes?

Enter Monte Carlo

- * Monte Carlo allows for solutions to complicated problems to be approximated by the use of random numbers
- * Modern version was invented during the development of the atom bomb. ENIAC was coded to carry out Monte Carlo simulations by John von Neumann.
- * Useful when closed-form expression or deterministic algorithm is not possible (fluids, computational biology, quantum physics)

Approach and equations

- * Randomly throw darts (in silico!) on dart board millions of times
- * Use 1x1 unit dart board (quarter of a circle)
- * Euclid: the area of a circle is proportional to r^2 and some constant.
- * Area of circle: $\pi \cdot r^2$
- * Probability: (favorable outcomes / total outcomes)
- * Probability of hitting in quarter of the unit circle: $\pi / 4$
- * Estimate probability: (random darts landed in circle / total random darts thrown)
- * Random darts: randomly choose x & y coordinate within unit circle. $A^2 + B^2 = C^2$ (Pythagorean theorem)